

Inheritance and Dynamic Binding

Laurent Mathy

Object-Oriented Programming Projects

April 17, 2020

Extending our student problem

- This time, we will assume students can take the course for both undergraduate and graduate credits.
- The grad students must do extra work: they write a thesis (report).
- We'd like our previous solution to the grading problem to continue to work.

Inheritance

- The record for graduate credit is the same as for undergraduate credit, except for additional properties related to the thesis.
- When you can think of a class as being like another class but for some extensions, you have a natural place for **inheritance**.
- In fact, inheritance can also be used for expressing exceptions (see your OO design classes).
- We will design 2 classes: one to represent undergrads (Core), and one to represent grads (called Grad).
- We'll use `Student_info` to represent any kind of students.

Core class

Core is similar to Student_info from previous lecture.

We add a **private** utility function to read the portion of student record that all students have in common.

```
1  class Core {
2  public:
3      Core();
4      Core(std::istream& is);
5      std::string name() const;
6      std::istream& read(std::istream&);
7      double grade() const;
8  private:
9      std::istream& read_common(std::istream&);
10     std::string n;
11     double midterm, final;
12     std::vector<double> homework;
13 };
```

Grad class

```
1  class Grad: public Core {
2  public:
3      Grad();
4      Grad(std::istream&);
5      double grade() const;
6      std::istream& read(std::istream&);
7  private:
8      double thesis;
9  };
```

- Grad **inherits** from Core, or Core is a **base class** of Grad.
- **public** in **public** Core means that the Core public interface is part of the Grad public interface.
 - **public** members of Core are also **public** members of Grad.
 - e.g. can call name() member function on a Grad.
- Grad objects will have 5 data members:
 - 4 inherited from Core + thesis.
- Grad will have 2 constructors and 4 member functions:
 - inherited name() and read_common();
 - **overridden** grade() and read().

Protection revisited

- Right now, all four data members and `read_common` in `Core` are inaccessible to member functions of `Grad`:
 - **private** members of a class are only accessible **to the class itself**, and its **friends**.
 - Use **protected** label to grant access to derived classes:
 - **protected** members are still accessible by the class itself and its friends, but also by **derived classes**.
-

```
1  class Core {
2  public:
3      Core();
4      Core(std::istream& is);
5      std::string name() const;
6      std::istream& read(std::istream&);
7      double grade() const;
8  protected: // Accessible to derived classes
9      std::istream& read_common(std::istream&);
10     std::string _name;
11     double midterm, final;
12     std::vector<double> homeworks;
13 };
```

Operations: Core

```
12  string Core::name() const { return _name; }
13
14  double Core::grade() const {
15      return ::grade(midterm, final, homeworks);
16  }
17
18  istream& Core::read_common(istream& in) {
19      // Read and store student's name and exam grades
20      in >> _name >> midterm >> final;
21      return in;
22  }
23
24  istream& Core::read(istream& in) {
25      read_common(in);
26      read_hws(in, homeworks);
27      return in;
28  }
```

Operations: Grad

```
11  istream& Grad::read(istream& in) {
12      read_common(in);
13      in >> thesis;
14      read_hws(in, homeworks);
15      return in;
16  }
17
18  double Grad::grade() const {
19      return std::min(Core::grade(), thesis);
20  }
```

- You **must** write `Core::grade()`, otherwise you get a recursive call to `Grad::grade()`.
- You *could* write `Core::read_common` and `Core::homework`, although these *are* members of `Grad`.

Inheritance and Constructors

Construction of derived objects:

- 1 Allocation of space for entire object.
- 2 Calling base-class constructor to initialise base-class part.
- 3 Initialisation of members of derived-class (via constructor initialisers).
- 4 Execution of derived-class constructor body, if any.

Use the constructor initialiser to specify the base-class constructor you want.

Initialiser names its base-class followed by a (possibly empty) list of arguments.

If no base-class constructor specified, then the default base-class constructor is run.

Constructors

```
1  class Core {
2  public:
3      // Default constructor for Core
4      Core(): midterm(0), final(0) { }
5      Core(std::istream& is) { read(is); }
6      // ...
7  };
8
9  class Grad: public Core {
10 public:
11     // Both constructors implicitly use Core::Core()
12     Grad(): thesis(0) { }
13     Grad(std::istream& is) { read(is); }
14     //...
15 };
```

Note that there is no requirement that the derived-class constructor take the same arguments as the base-class constructors.

Polymorphism

We had a non-member compare function used by sort to sort student records by name:

```
35 bool compare(const Core& c1, const Core& c2) {  
36     return c1.name() < c2.name();  
37 }
```

We can use this code to compare two Core objects, two Grad objects, or even a Core and a Grad.

```
22 Grad g1(in);  
23 Grad g2(in);  
24  
25 Core c1(in);  
26 Core c2(in);  
27  
28 compare(g1, g2);  
29 compare(c1, c2);  
30 compare(c1, g1);
```

The reason why this works is because every Grad has a Core part.

Polymorphism (2)

Because Grad **inherits** from Core, you can use a Grad where a Core is expected.

- A reference parameter to a Core will refer to the Core portion of a Grad.
- A pointer parameter to a Core will point to the Core portion of a Grad if a pointer to Grad passed instead (Grad* converted to Core* by compiler).
- Object of type Core corresponds to Core portion of Grad if a Grad object is assigned/passed instead.

We say that Grad is a **subtype** of Core, noted $\text{Grad} <: \text{Core}$.

Polymorphism(3): a new compare function

Suppose that instead of sorting students by name, we want to sort them by final grade:

```
10 bool compare_grades(const Core& c1, const Core& c2) {  
11     return c1.grade() < c2.grade();  
12 }
```

Polymorphism(3): a new compare function

Suppose that instead of sorting students by name, we want to sort them by final grade:

```
10 bool compare_grades(const Core& c1, const Core& c2) {  
11     return c1.grade() < c2.grade();  
12 }
```

As Grad redefines the grade function, this compare_grades functions sometimes gives the wrong answer, because it **always** invokes the Core::grade function, as c1 and c2 *are* references to Core objects.

We need a way for the compare_grades function to invoke the right grade function, based on the **actual** type of the object that we pass:

- If c1 or c2 are Grads, we want Grad::grade.
- If c1 or c2 are Cores, we want Core::grade.
- This must be done at **run-time**, as the **dynamic** types of the passed objects can only be determined at run-time.

Polymorphism (4): **virtual**

To support this:

```
1  class Core {  
2  public:  
3      virtual double grade() const; // added `virtual`  
4      // ...  
5  };
```

- grade is now a **virtual** function.
- **virtual** keyword may be used *only* in class definitions: do not repeat it in function definitions.
- **virtual** is inherited, so no need to repeat it in Grad, though doing it doesn't hurt.

Note to Java programmers: in Java, all member functions are **virtual** by default. In C++, you must turn this **dynamic binding** on explicitly!

Dynamic Binding

The run-time selection of the **virtual** function is relevant only for references and pointers.

If a function is called on behalf of an object, you get the version of the function corresponding to the object type. In other words, the type of an object is immutable!

```
1 // Incorrect implementation
2 bool compare_grades(Core c1, Core c2) {
3     return c1.grade() < c2.grade();
4 }
```

- c1 and c2 are always Core.
- If you pass a Grad to this function, only the Core part gets copied.
- Because we said the parameters are Core objects, the calls to grade are **statically** bound at **compile-time** to Core::grade.

Dynamic-binding only applies to references and pointers.

Dynamic Binding (2)

```
39 Core c;  
40 Grad g;  
41 Core* p = &c;  
42 Core& r = g;  
43  
44 c.grade(); // Static binding to Core::grade()  
45 g.grade(); // Static binding to Grad::grade()  
46 p->grade(); // Dynamic binding to type object p points to  
47 r.grade(); // Dynamic binding to type object r refers to
```

Polymorphic call:

The type of the reference or pointer is fixed, but the type of the object referred or pointed to can be the corresponding type or any type derived from it.

virtual and pure virtual

- Non-virtual functions can be declared, without being defined, if they are not called.
- **virtual** functions *must* be defined, whether they are called or not.
 - You'll get weird compile errors if not.
- If there is no meaningful implementation for a virtual function, make it a **pure virtual** function.

```
1  class Abstract {  
2  public:  
3      virtual int pure() const = 0; // Pure virtual  
4      // ...  
5  };
```

Such (abstract) classes cannot be instantiated: they can only serve as base for derived classes.

A program dealing with only undergrad records

```
10  vector<Core> students;
11  Core record;
12  string::size_type max_len = 0;
13  // Read and store the student records
14  while (record.read(cin)) { // Core::read()
15      max_len = max(max_len, record.name().size());
16      students.push_back(record);
17  }
18  // Alphabetize the student records
19  sort(students.begin(), students.end(), compare);
20  // Write the names and grades
21  for (auto& s : students) {
22      cout << s.name()
23           << string(max_len + 1 - s.name().size(), ' ');
24      try {
25          cout << s.grade() << endl; // Core::grade
26      } catch (domain_error e) {
27          cerr << e.what() << endl;
28      }
29  }
```

A program dealing with only grad records

```
10  vector<Grad> students; // Different type in vector
11  Grad record; // Different type into which to read
12  string::size_type max_len = 0;
13  // Read and store the student records
14  while (record.read(cin)) { // Grad::read()
15      max_len = max(max_len, record.name().size());
16      students.push_back(record);
17  }
18  // Alphabetize the student records
19  sort(students.begin(), students.end(), compare);
20  // Write the names and grades
21  for (auto& s : students) {
22      cout << s.name()
23           << string(max_len + 1 - s.name().size(), ' ');
24      try {
25          cout << s.grade() << endl; // Grad::grade
26      } catch (domain_error e) {
27          cerr << e.what() << endl;
28      }
29  }
```

Towards a program that deals with both types of records

We need to eliminate the following type dependencies:

- Definition of the vector.
- Definition of the local variable.
- Calling the right read function.
- Calling the right grade function.

As read and grade are **virtual**, last two points have been solved.

The only problem is that our code makes **statically-bound** calls to these functions. We need to turn **dynamic** binding on.

We can use `Core*` where we used `Core` or `Grad`, and let our users allocate memory.

We'll also need a function to compare Cores identified by pointers:

```
10 bool compare_Core_ptrs(const Core* cp1, const Core* cp2) {  
11     return compare(*cp1, *cp2);  
12 }
```

Dealing with both types of records (2)

```
15  vector<Core*> students; // Store pointers, not objects
16  Core* record; // Temporary must be a pointer as well
17  char ch;
18  string::size_type max_len = 0;
19
20  // Read and store the student records
21  while (cin >> ch) {
22      if (ch == 'U')
23          record = new Core; // Allocate a Core object
24      else
25          record = new Grad; // Allocate a Grad object
26      record->read(cin); // Virtual call
27      max_len = max(max_len, record->name().size());
28                      // ^ Dereference
29      students.push_back(record);
30  }
31
32  // Pass the version of compare() that works on pointers
33  sort(students.begin(), students.end(), compare_Core_ptrs);
```

Dealing with both types of records (3)

```
35  // Write the names and grades
36  for (auto s : students) {
37      // s is a pointer
38      cout << s->name() // Dereference to call function
39          << string(max_len + 1 - s->name().size(), ' ');
40      try {
41          cout << s->grade() << endl; // Dereference to call
42
43      } catch (domain_error e) {
44          cerr << e.what() << endl;
45      }
46      delete s; // Free the object allocated when reading
47  }
```

Virtual destructors

Our previous program *almost* works:

- When we **delete** the records, it is always through a `Core*`.
- But these `Core*` can point to either `Core` or `Grad` objects.
- Which destructor to call and how much space to reclaim?
 - Sounds like exactly what the **virtual** mechanism handles.

virtual destructor:

```
1  class Core {  
2  public:  
3      virtual ~Core() { }  
4      // ...  
5  };
```

- A **virtual** destructor is needed any time a derived type object can be destroyed through a pointer to base.
- If no other reason to have destructor, then that destructor has no work to do and is empty
- No need to add destructor to `Grad`: virtual property of destructor is inherited and synthesised destructor is fine.

Another Solution

In previous solution, our users had to do memory management for us: this is messy and error prone.

We'll define a **handle** (a.k.a. **proxy**) class, based on a `Core*` that does the memory bookkeeping itself.

Student_info handle class

```
10  class Student_info {
11  public:
12      // Constructors and copy control
13      Student_info() { }
14      Student_info(std::istream& is) { read(is); }
15      Student_info(const Student_info&);
16      Student_info& operator=(const Student_info&);
17      // Operations
18      std::istream& read(std::istream&);
19      std::string name() const {
20          if (cp) return cp->name();
21          else throw std::runtime_error("uninitialized Student");
22      }
23      double grade() const {
24          if (cp) return cp->grade();
25          else throw std::runtime_error("uninitialized Student");
26      }
27      static bool compare(const Student_info& s1, const Student_info& s2) {
28          return s1.name() < s2.name();
29      }
30  private:
31      std::unique_ptr<Core> cp;
32  };
```

Student_info handle class (2)

- As Core has a virtual destructor, the Student_info destructor will work properly, whether it represents a Core or Grad
- Student_info::read will allocate the appropriate space.
- As grade is **virtual**, we will get correct version as called through Core* pointer.
- compare has been made a **static function**, it:
 - is associated with class, not with any particular object;
 - cannot access non-**static** members;
 - call Student_info::compare.

Reading the handle

```
7  istream& Student_info::read(istream& is) {
8      char ch;
9      is >> ch; // Get record type
10
11     // Assignment to `cp` will free if needed
12     if (ch == 'U')
13         cp = std::make_unique<Core>(is);
14     else
15         cp = std::make_unique<Grad>(is);
16
17     return is;
18 }
```

Copying handle objects

We need a copy constructor and assignment operators to manage the Core*.

Question: when we copy, are we copying a Core or a Grad? There is no easy way to know!

We solve this problem by giving Core and its derived classes a new **virtual** function:

```
1  class Core {
2      friend class Student_info;
3  protected:
4      virtual std::unique_ptr<Clone> clone() const
5          { return std::make_unique<Core>(*this); }
6      // ...
7  };
8  class Grad {
9  protected:
10     virtual std::unique_ptr<Clone> clone() const
11         { return std::make_unique<Grad>(*this); }
12     // ...
13 };
```

These **virtual** clone functions call the synthesised copy constructor for Core and Grad.

Copying Handle objects (2)

- We do not want `clone` as member of `public` interface, so it is made `protected`.
- Because `clone` is `protected`, we make `Student_info` `class` a `friend` of `Core`.
 - All member functions of `Student_info` are now `friends` with `Core`.
- When derived class redefines a function from base class, it usually does it **exactly**: parameter list and return type are identical.
 - However, if base-class function returns a pointer or reference to the base class, then the derived class can return a pointer or reference to the derived class.
- `friendship` is not inherited, but no need to make `Student_info` a `friend` of `Grad`.
 - `Student_info` never refers to `Grad::clone` directly,
 - only through `virtual` calls to `Core::clone`.

Copying and assignement

```
20 Student_info::Student_info(const Student_info& s) {
21     if (s.cp)
22         cp = s.cp->clone();
23 }
24
25 Student_info& Student_info::operator=(const Student_info& s) {
26     if (&s != this) {
27         if (s.cp)
28             cp = s.cp->clone();
29         else
30             cp = nullptr;
31     }
32
33     return *this;
34 }
```

Using the handle class

```
10  vector<Student_info> students;
11  Student_info record;
12  string::size_type max_len = 0;
13  // Read and store the student records
14  while (record.read(cin)) {
15      max_len = max(max_len, record.name().size());
16      students.push_back(record);
17  }
18  // Alphabetize the student records
19  sort(students.begin(), students.end(), Student_info::compare);
20  // Write the names and grades
21  for (auto& s : students) {
22      cout << s.name()
23           << string(max_len + 1 - s.name().size(), ' ');
24      try {
25          cout << s.grade() << endl;
26      } catch (domain_error e) {
27          cerr << e.what() << endl;
28      }
29  }
```


Subtleties: inheritance and containers

```
1 vector<Core> students;  
2 Grad g(cin);  
3 students.push_back(g);
```

What happens?

Subtleties: inheritance and containers

```
1 vector<Core> students;  
2 Grad g(cin);  
3 students.push_back(g);
```

What happens?

We are allowed to store a Grad in a vector<Core>.
But only the Core part of the Grad will be stored!

Subtleties: functions

If base-class and derived-class have a function with the same name but different signatures, they behave as unrelated functions.

```
1 void Core::regrade(double d) { final = d; }
2 void Grad::regrade(double d1, double d2) {
3     final = d1;
4     thesis = d2;
5 }
```

If `r` is a reference to `Core`

```
1 r.regrade(100); // OK, call Core::regrade
2 r.regrade(100, 100);
3 // Compile error: Core::regrade takes 1 argument
```

This second call is an error even if `r` actually refers to a `Grad`!

Subtleties: functions (2)

If `r` is a reference to `Grad`

```
1 r.regrade(100); // Comp. error: Grad::regrade takes 2 arguments
2 r.regrade(100, 100); // OK, call Grad::regrade
```

Even though there is a base-class version that takes a single argument, it is effectively hidden by the existence of `regrade` in the derived class.

If you really want the base-class version, you need:

```
r.Core::regrade(100);
```

To use `regrade` as a **virtual** function, it must have the same interface in both the base and derived class:

```
1 virtual void Core::regrade(double d, double = 0) { final = d; }
2 void Grad::regrade(double d1, double d2) {
3     final = d1;
4     thesis = d2;
5 }
```
